




LEADS 3-Day Camp

Session 3: Data Management: Relational Databases and SQL

Il-Yeol Song, Ph.D.
Professor
College of Computing & Informatics
Drexel University
Philadelphia, PA 19104
Song@drexel.edu
<http://www.cci.drexel.edu/faculty/song/>

INFO 607, Il-Yeol Song

Hello!



INFO 607, Il-Yeol Song

Instructor

- Professor, College of Computing & Informatics, Drexel University,
- PhD in CS, LSU, Baton Rouge, USA, 1988
- CCI PhD Program Director (2010-2015)
- CVDI Deputy Director (2012-2014)
- Research Topics: *Conceptual Modeling, Data Warehousing, Big Data Management and Data Analytics, Smart Aging*
- Elected as an **ER Fellow**, 2012
- Named an **ACM Distinguished Scientist** in 2013
- Received **Peter Chen Award in Conceptual Modeling** in 2015
- Four teaching awards from Drexel (1991, 2000, 2001, 2011) including **Lindback Distinguished Teaching Award** (2001)
- **Co-Editor-in-Chief**, *Journal of Computing Science & Engineering*
- **Consulting Editor**, *Data & Knowledge Engineering*
- **Co-Chair**, iSchool Model Data Science Curriculum Committee
- **Chair**, IEEE Big Data and Smart Computing (BigComp) Conference
- Published about 200+ papers

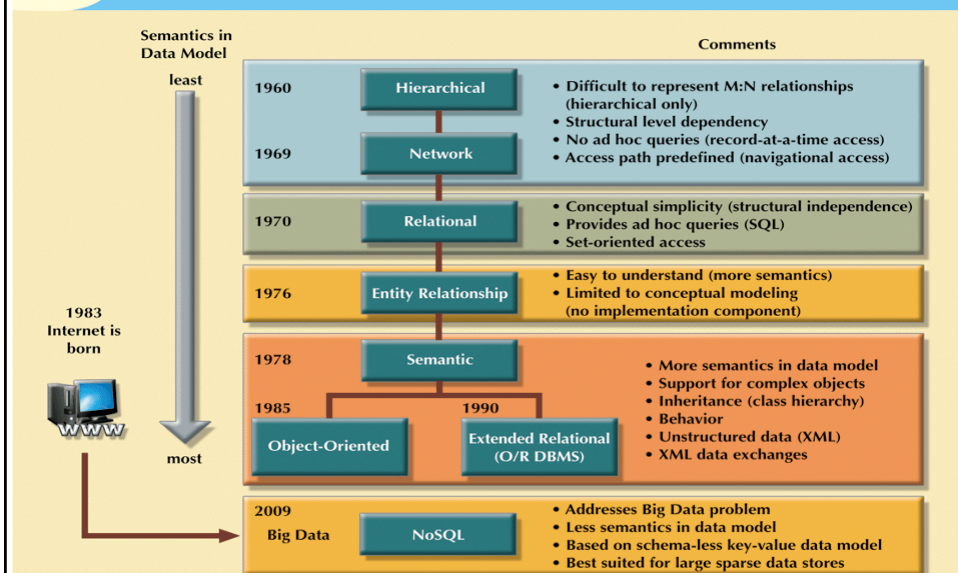
Acknowledgement

- Materials from:
 - Lecture Notes from Drexel CCI INFO 605 and 606 Class (Il-Yeol Song)
 - Some diagrams and examples from
Carlos Coronel and Steven Morris, Database Systems: Design, Implementation, and Management, 13th Edition, 2017. Cengage Learning.

Contents

- Evolution of Data Models
- Relational Databases
- SQL
 - Basic SQL commands
 - SQL for Data Engineering
 - SQL for Analytics
- NoSQL Databases
- MongoDB Basic CRUD Operations

FIGURE 2.6 The evolution of data models



SOURCE: Course Technology/Cengage Learning

Basics in the Relational Model

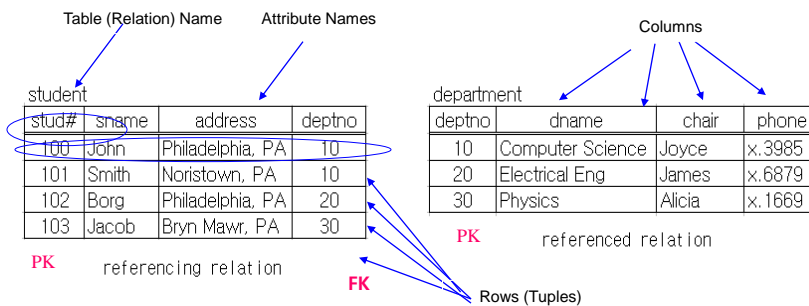
- Data representation: a set of tables with structured data
- Primary keys (Candidate key, Alternate keys)
- Foreign Key
- Referential Integrity constrains
- Integrity constraints in RDBMSs (PK, RI, CHECK, Triggers)
- Null values
- Indexes
- Metadata and Data Dictionary
- Relational schema design using ER diagrams
- Normalization (FD, MVD, JD) and Normal Forms
- SQL (*Aggregation, Subquery, JOIN*, PL/SQL functions/procedures/triggers)
- ACID property in multi-user *transactions*

Terminology in RDB

Relational Schema

STUDENT (stud#, sname, address, deptno)

DEPARTMENT (deptno, dname, chair, phone)



A **Primary Key (PK)** uniquely identifies each row.

A **Foreign Key (FK)** is a set of attributes that are primary key values from another table.

Characteristics of Relational Model




TABLE 3.1 Characteristics of a Relational Table

1	A table is perceived as a two-dimensional structure composed of rows and columns.
2	Each table row (tuple) represents a single entity occurrence within the entity set.
3	Each table column represents an attribute, and each column has a distinct name.
4	Each row/column intersection represents a single data value.
5	All values in a column must conform to the same data format.
6	Each column has a specific range of values known as the attribute domain .
7	The order of the rows and columns is immaterial to the DBMS.
8	Each table must have an attribute or a combination of attributes that uniquely identifies each row.

Example of a DDL in SQL Defining a Relational Schema

student	deptno	dname	chair	phone
stud#	sname	address	deptno	
100	John	Philadelphia, PA	10	
101	Smith	Noristown, PA	10	
102	Borg	Philadelphia, PA	20	
103	Jacob	Bryn Mawr, PA	30	

referencing relation

deptno	dname	chair	phone
10	Computer Science	Joyce	x.3985
20	Electrical Eng	James	x.6879
30	Physics	Alicia	x.1669

referenced relation

```
CREATE TABLE Department (
    DeptNo      NUMBER(5)    PRIMARY KEY,
    DName       VARCHAR2(15) NOT NULL,
    Chair       VARCHAR2(20),
    Phone       CHAR(10),
);
```


```
CREATE TABLE Student (
    Stud#       NUMBER(8),
    SName       VARCHAR2(20) NOT NULL,
    Address     VARCHAR2(40),
    DeptNo      NUMBER(5),
    CONSTRAINT Stud_PK PRIMARY KEY (Stud#),
    CONSTRAINT Stud_FK FOREIGN KEY (DeptNo)
    REFERENCES Department (DeptNo)
);
```

ER Model: Popular Notations


Chen




Visio 's Relational notation




UML




IDEF1X



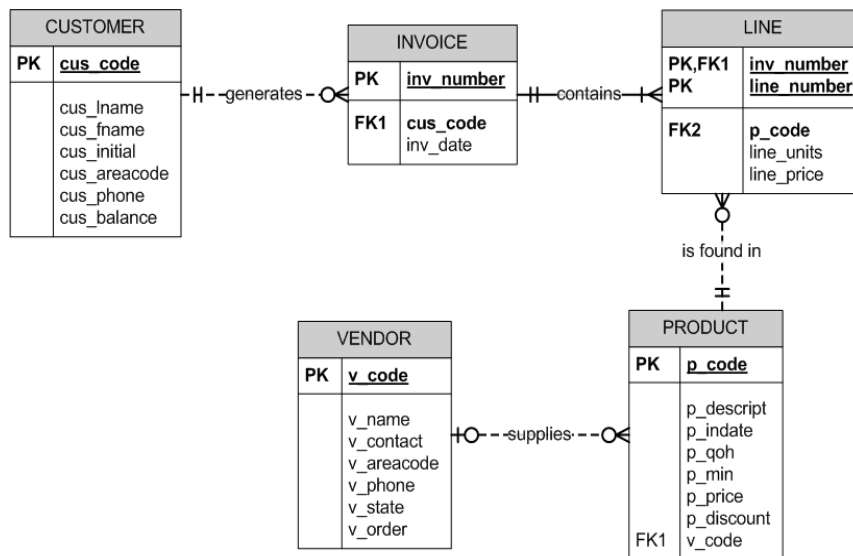
Crow's foot with Relational Notation in Visio



Visio IDEF1X with Crow's foot




ERD for Relational DBs



Why Relational Databases?

- Logical simplicity of the schema: **Tables**
- Easy, powerful, standard database language: **SQL**
- Transaction reliability: **ACID property** (Atomicity, Consistency, Isolation, Durability)
- Ad-hoc query processing
- Mature and reliable technologies
- Commercial investment for the last 40 years
- A large group of man-power and user bases

2020 Kaggle Data Science & Machine Learning Survey



<https://www.kaggle.com/code/paultimothymooney/2020-kaggle-data-science-machine-learning-survey/notebook>

Top Desired Skills in Data Science (2021)

11 Skills Data Scientists Need

1. Data Visualization
2. Python
3. SQL/NoSQL
4. Social Media Mining
5. Fundamental Statistics
6. Natural Language Processing/Machine Learning
7. Microsoft Excel
8. High-Level Math
9. Teamwork
10. Communication
11. Business Savvy

Source: <https://bootcamp.berkeley.edu/blog/data-scientist-skills/>

Introduction to SQL

- Pronounced ‘see-quel’
- ANSI prescribes a standard SQL
- Developed by IBM in 1974, System R project
- First commercial implementation by Oracle (1976)
- ANSI/ISO Standard Versions:
 - SQL 86 (SQL 1)
 - SQL 89 (SQL 1 Level 2): RI integrity
 - SQL 92 (SQL 2)
 - SQL:1999 for OO features
 - SQL:2003 (XML features)
 - SQL: 2011 (Temporal data, windows functions)
 - SQL: 2016: JSON
 - SQL: 2019: Multi-dimensional Array (MDArray type and operations)
- Several dialects and extensions exist

SQL Topics

- **SQL Basics**
 - Basic DDL and DML
 - Simple commands on a Single table
 - Joins (inner join, outer join, natural joins, multi-table joins)
 - Subquery
- **SQL for Data Engineering**
 - Date manipulation and Date Arithmetic
 - String manipulation
 - Merge
- **SQL for Analytics**
 - Aggregation, grouping, and ordering
 - OLAP functions
 - Ranking, partition by, buckets
 - Creating new derived columns

Two major Components of SQL

Data Definition Language (**DDL**)

- CREATE TABLE
- ALTER TABLE
- DROP TABLE
- CREATE/DROP VIEW
- CREATE/DROP INDEX

Data Manipulation Language(**DML**)

- UPDATE TABLE
- INSERT INTO
- DELETE FROM
- **SELECT** FROM

DDL manages metadata
DML manages data

A Typical SQL Statement and Rules

STUDENT

ID	Name	EnrollYear

- **SELECT** enrollyear **FROM** student **WHERE** name = 'John';
- SQL statement comprises
 - **Keywords**
 - **Case-insensitive**
 - Basic command set: fewer than 100 English words
 - Space-independent
 - **User-defined word** (table names, attribute names)
 - **Case-insensitive**
 - **Data**
 - **Case-sensitive** (In MS Access, case-insensitive)
 - Ends with a semi-colon

An Overview of SQL Commands

DDL

--Changing attribute type

```
ALTER TABLE project
  MODIFY (budget
    NUMBER(9,2));
```

--Adding an attribute

```
ALTER TABLE project ADD
  (manager CHAR(10) );
```

-- Delete a column from a table

```
ALTER TABLE project
  DROP COLUMN
  manager;
```

--Removing a table from database

```
DROP TABLE project;
```

DML

--Inserting a row to a table

```
INSERT INTO project VALUES (1234,
  'Perfect Project', NULL, 'John');
```

--Changing a value of attribute

```
UPDATE project SET budget =
  1.1*budget WHERE projno > 1000;
```

--Deleting a row from a table

```
DELETE FROM project WHERE
  manager = 'John';
```

Queries using SELECT Statement

SELECT	Specifies which columns are to appear in output
FROM	Specifies table(s) to be used
WHERE	Filters rows with conditions
GROUP BY	Forms groups of rows with the same column value.
HAVING	Filters groups subject to some condition.
ORDER BY	Specifies the order of the output.

- *Only SELECT and FROM are mandatory.*
- *Order of the clauses cannot be changed.*

An Overview of Simple Commands

- | | |
|--|--|
| (1) SELECT * FROM
Instructor; | (6) SELECT * FROM Instructor
WHERE deptCode IS NULL ; |
| (2) SELECT DISTINCT fName
FROM Instructor; | (7) SELECT * FROM Instructor
WHERE deptCode IS NOT NULL ; |
| (3) SELECT fName, lName, ssn
FROM Instructor
WHERE deptCode = 'math'; | (8) SELECT * FROM Instructor
WHERE bonus BETWEEN 500 AND 1000; |
| (4) SELECT fName, lName
FROM Instructor
WHERE bonus > 1000; | (9) SELECT * FROM HR
WHERE hireDate
BETWEEN '01-MAY-2022' AND '31-MAY-2022'; |
| (5) SELECT * FROM Instructor
WHERE (bonus >= 500 AND bonus <= 1000); | (10) SELECT * FROM Instructor
WHERE bonus IN (100, 200, 300); |

An Overview of Simple Commands

- | | |
|--|---|
| <p>(11) SELECT * FROM
Instructor
WHERE position = 'assistant';</p> <p>(12) SELECT * FROM R
WHERE lName > 'S';</p> <p>(13) SELECT * FROM
Instructor
WHERE fName LIKE 'J%';</p> <p>(14) SELECT * FROM
Instructor
WHERE fName LIKE 'J_h_';</p> | <p>(15) SELECT fName, lName FROM
Instructor
WHERE address LIKE '%Houston,TX%';</p> <p>(16) SELECT (salary+bonus) FROM
Instructor;</p> <p>(17) SELECT (salary+bonus) AS
total_income FROM Instructor;</p> <p>(18) SELECT ((A1/3.14)*2.54) AS
A1_IN_CM FROM R;</p> <p>(19) SELECT SYSDATE AS TODAY
FROM DUAL;</p> |
|--|---|
- (20) SELECT FName, LName, Bday,
TRUNC (MONTHS_BETWEEN (SYSDATE, Bday)/12) AS "Actual Age"
FROM Person;

Join Operation

- JOIN is the relational operation that combines the data stored in two tables, creating a new result table
- JOIN is the most important operation in relational query processing
- **JOIN is performed between two tables that have a PK-FK relationship.**
 - If your FROM clause contains **2 tables**, you need **one join**.
 - If your FROM clause contains **3 tables**, you need **two joins**.
 - If your FROM clause contains **N tables**, you need **N-1 joins**.
- **Joining between two tables that do not have a PK-FK relationship may NOT BE CORRECT.**

JOIN: Two Syntax


JOIN is usually done through **PK-FK chains**

- For each product, find their vendor code and names

SELECT P.P_Code, V.V_Code, V.V_Name

FROM Product **P**, Vendor **V**

WHERE **P.V_Code = V.V_code;**



SELECT P.P_Code, V.V_Code, V.V_Name

FROM Product **P INNER JOIN** Vendor **V**

ON **P.V_Code = V.V_code;**

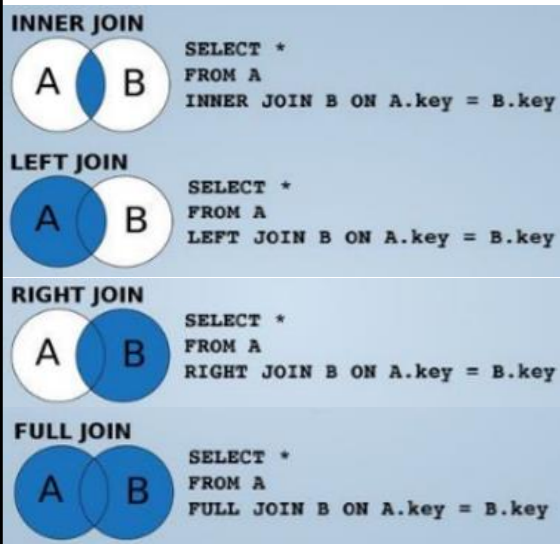
Types of JOINS

TABLE 8.1

SQL JOIN EXPRESSION STYLES

JOIN CLASSIFICATION	JOIN TYPE	SQL SYNTAX EXAMPLE	DESCRIPTION
CROSS	CROSS JOIN	SELECT * FROM T1, T2	Returns the Cartesian product of T1 and T2 (old style)
		SELECT * FROM T1 CROSS JOIN T2	Returns the Cartesian product of T1 and T2
INNER	Old-style JOIN	SELECT * FROM T1, T2 WHERE T1.C1=T2.C1	Returns only the rows that meet the join condition in the WHERE clause (old style); only rows with matching values are selected
	NATURAL JOIN	SELECT * FROM T1 NATURAL JOIN T2	Returns only the rows with matching values in the matching columns; the matching columns must have the same names and similar data types
	JOIN USING	SELECT * FROM T1 JOIN T2 USING (C1)	Returns only the rows with matching values in the columns indicated in the USING clause
OUTER	JOIN ON	SELECT * FROM T1 JOIN T2 ON T1.C1=T2.C1	Returns only the rows that meet the join condition indicated in the ON clause
	LEFT JOIN	SELECT * FROM T1 LEFT OUTER JOIN T2 ON T1.C1=T2.C1	Returns rows with matching values and includes all rows from the left table (T1) with unmatched values
	RIGHT JOIN	SELECT * FROM T1 RIGHT OUTER JOIN T2 ON T1.C1=T2.C1	Returns rows with matching values and includes all rows from the right table (T2) with unmatched values
	FULL JOIN	SELECT * FROM T1 FULL OUTER JOIN T2 ON T1.C1=T2.C1	Returns rows with matching values and includes all rows from both tables (T1 and T2) with unmatched values

Types of JOINS



NATURAL JOIN:

```
SELECT *
FROM table_A
NATURAL JOIN table_B;
```


SELF JOIN:

```
SELECT *
FROM table_A X, table_A Y
WHERE X.A = Y.A;
```

CROSS JOIN:


```
SELECT *
FROM table_A
CROSS JOIN table_B;
```

Types of JOINS in Tableau



INNER JOIN

- For each invoice, find the product description, #units, and price of items in the invoice.



SQL Plus

```
SQL> SELECT INVOICE.INV_NUMBER, PRODUCT.P_CODE, P_DESCRIPTION, LINE_UNITS, LINE_PRICE
2 FROM INVOICE JOIN LINE ON INVOICE.INV_NUMBER = LINE.INV_NUMBER
3 JOIN PRODUCT ON LINE.P_CODE = PRODUCT.P_CODE;
```


INV_NUMBER	P_CODE	P_DESCRIPTION	LINE_UNITS	LINE_PRICE
1001	13-Q2/P2	7.25-in. pwr. saw blade	1	14.99
1001	23109-HB	Claw hammer	1	9.95
1002	54778-2T	Rat-tail file, 1/8-in. fine	2	4.99
1003	2238/QPD	B&D cordless drill, 1/2-in.	1	38.95
1003	1546-QQ2	Hrd. cloth, 1/4-in., 2x50	1	39.95
1003	13-Q2/P2	7.25-in. pwr. saw blade	5	14.99
1004	54778-2T	Rat-tail file, 1/8-in. fine	3	4.99
1004	23109-HB	Claw hammer	2	9.95
1005	PVC23DRT	PVC pipe, 3.5-in., 8-ft	12	5.87
1006	SM-18277	1.25-in. metal screw, 25	3	6.99
1006	2232/QTY	B&D jigsaw, 12-in. blade	1	109.92
1006	23109-HB	Claw hammer	1	9.95
1006	89-WRE-Q	Hicut chain saw, 16 in.	1	256.99
1007	13-Q2/P2	7.25-in. pwr. saw blade	2	14.99
1007	54778-2T	Rat-tail file, 1/8-in. fine	1	4.99
1008	PVC23DRT	PVC pipe, 3.5-in., 8-ft	5	5.87
1008	WR3/TT3	Steel matting, 4'x8'x1/6", .5" mesh	3	119.95
1008	23109-HB	Claw hammer	1	9.95

18 rows selected.

SQL> _

LEFT OUTER JOIN

- *For ALL the Vendors*, find all the products they provide. Include all the vendors whether or not they provide any product.



SQL Plus

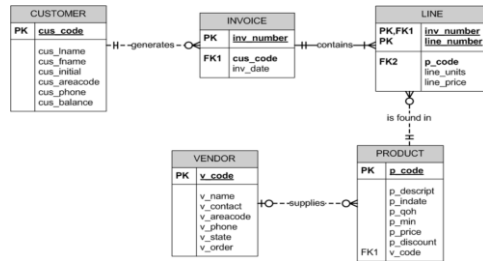
```
SQL> SELECT P_CODE, VENDOR.V_CODE, V_NAME
2 FROM VENDOR LEFT JOIN PRODUCT ON VENDOR.V_CODE = PRODUCT.V_CODE;
```

P_CODE	V_CODE	V_NAME
110ER/31	25595	Rubicon Systems
13-Q2/P2	21344	Gomez Bros.
14-Q1/L3	21344	Gomez Bros.
1546-QQ2	23119	Randsets Ltd.
1558-QW1	23119	Randsets Ltd.
2232/QTY	24288	ORDVA, Inc.
2232/QWE	24288	ORDVA, Inc.
2238/QPD	25595	Rubicon Systems
23109-HB	21225	Bryson, Inc.
54778-2T	21344	Gomez Bros.
89-WRE-Q	24288	ORDVA, Inc.
SM-18277	21225	Bryson, Inc.
SW-23116	21231	D&E Supply
WR3/TT3	25595	Rubicon Systems
	22567	Dome Supply
	21226	Superloo, Inc.
	24004	Brackman Bros.
	25501	Damal Supplies
	25443	B&K, Inc.

19 rows selected.

SQL> _

Multi-table joins



54

List customer last name, invoice number, invoice date, and product description for all invoices for customer 10014.

```

SELECT      CUS_LNAME, I.INV_NUMBER, INV_DATE, P_DESCRIPT
FROM        CUSTOMER C, INVOICE I, LINE L, PRODUCT P
WHERE        C.CUS_CODE = I.CUS_CODE
              AND I.INV_NUMBER = L.INV_NUMBER
              AND L.P_CODE = P.P_CODE
              AND C.CUS_CODE = 10014

ORDER BY    INV_NUMBER;
  
```

Subqueries

- Subquery (or nested query): a complete **SELECT** statement embedded within another **SELECT** statement.

```

SELECT      select_list
FROM        table
WHERE        expr operator
              (SELECT select_list
               FROM      table);
  
```

- The results of this *inner* **SELECT** statement are used in the *outer* **SELECT** statement to help determine the final result.
- A subquery produces a *temporary table* with results
- Can be used in **SELECT**, **FROM**, **WHERE**, and **HAVING** clauses of an outer **SELECT** statement and also in **INSERT**, **UPDATE**, **DELETE** statements


Subqueries

- Find all vendor names which supplies USB

```
SELECT v_name, v_phone
FROM Vendor
WHERE v_code IN (SELECT v_code
                  FROM Product
                  WHERE p_descript = 'USB');
```

- Can we use “=” instead of “IN”?

```
SELECT v_name, v_phone
FROM Vendor
WHERE v_code = (SELECT v_code
                 FROM Product
                 WHERE p_descript = 'USB');
```



SQL Topics

- SQL Basics**
 - Basic DDL and DML
 - Simple commands on a Single table
 - Joins (inner join, outer join, natural joins, multi-table joins)
 - Subquery
- SQL for Data Engineering**
 - Date manipulation and Date Arithmetic
 - String manipulation
 - Merge
- SQL for Analytics**
 - Aggregation, grouping, and ordering
 - OLAP functions
 - Ranking, partition by, buckets
 - Creating new derived columns

DATE Manipulation & Formats

- To insert a date using a default format:
`INSERT INTO Loan (Loan#, LoanDate) VALUES (12, '03-FEB-2019');`
- To insert a date using a non-default format:
`INSERT INTO Loan (Loan#, LoanDate)
VALUES (12, TO_DATE('06/02/2022', 'MM/DD/YYYY'));`
- Show today's date in the form of MM/DD/YYYY format
`SELECT SYSDATE, TO_CHAR(SYSDATE, 'MM/DD/YYYY') as
CurrDate FROM DUAL;`
- Show today's date including time
`SELECT SYSDATE, TO_CHAR(SYSDATE, 'yyyy-mm-dd hh24:mi:ss')
as CurrDateTime FROM DUAL;`

DATE Arithmetic

- Find products whose INDATE is more than 90 days old
`SELECT P_Code, P_Desc, P_Indate
FROM Product
WHERE P_Indate <= SYSDATE - 90;`
- Find all the orders received in Feb 2022.
`SELECT Order#, Odate
FROM ORDER
WHERE Odate BETWEEN '01-Feb-22' AND '28-Feb-22';`

DATE Extraction

- **EXTRACT** function: returns numerical value of YEAR, MONTH, or DATE

Employee (E#, Ename, DoB)

-- Find the years of employees were born

```
SELECT  Ename, EXTRACT (YEAR FROM DoB)
FROM    Employee;
```

-- Month and day can also be extracted

```
SELECT Ename, EXTRACT (MONTH FROM DoB),
        EXTRACT (DAY FROM DoB)
FROM Employee;
```

STRING Manipulation

- **Concatenating strings**
SQL> `SELECT Lname||', '||Fname "FULL NAME" FROM Emp;`
will display
FULL NAME
Bond, James
- **Handling Capitalization**
SQL> `SELECT Lname, INITCAP(Lname), UPPER(Lname), LOWER(Lname) FROM Emp;`
- **Capitalization:** Show only the first letter in capital and the rest in lower
SQL> `SELECT INITCAP(LOWER(Lname)) FROM Emp;`
- **Padding blanks to strings**
RPAD: Pad to the right side of the column with left justification
LPAD: Pad to the left side of the column with right justification
SQL> `SELECT RPAD(Lname, 15, '.'), Age, LPAD(Lname, 15), FROM EMP:`

Clinton.....	53Clinton
Gore.....	52Gore

STRING Manipulation

/* **TRIM** function is used to remove all leading or trailing characters (or both) from a character string.*/

```
SQL>SELECT TRIM(" LEADING Fellows ") FROM DUAL;
```

/* Use of **LTRIM**(left trim) and **RTRIM**(right trim) functions*/

```
SQL> SELECT Lname, LTRIM(Lname, 'SA'), RTRIM(Lname, 'S')
FROM Emp;
```

(LTRIM prints M from SAM, RTRIM prints WALE from WALES)

/* Remove **."** at the end and **"THE** from the front */

```
SQL> SELECT LTRIM ( RTRIM (Title,','), ' "THE ') FROM
Magazine;
```

Will convert

MY DARING."

"THE GOD FATHER."

into

MY DARING

GOD FATHER

STRING Manipulation

- **SUBSTR**

- Syntax : SUBSTR(char, m, n)
- Purpose : Returns a portion of the char, **beginning at character 'm' up to 'n' characters**. If 'n' is omitted, result is returned up to the end of char.

- Example : **SELECT SUBSTR(first_name,1,3)**
FROM employee;

Output :

SUB

JOH

KEV

JEA

LYN

LES

CYN

STRING Manipulation

/* SUBSTR(String, starting_position, #chars to be returned) */

/* Extract numeric part, add 1000, and concatenate */

```
SQL> SELECT Lname, E_ID,
        'S' || TO_CHAR( TO_NUMBER (SUBSTR (E_ID,2,3) ) + 1000)
        "NEW E_#"
        FROM Emp;
```

Will display

LNAME	E_ID	NEW E_#
Jones	E001	S1001
Wales	E002	S1002

REPLACE() and TRANSLATE()

<https://www.databasestar.com/oracle-translate/>

TRANSLATE(string, from_string, to_string): From the *string*, replace each character of *from_string* into the corresponding character of the *to_string*, *char by char (case sensitive)*

REPLACE (string, from_string, to_string): From the *string*, replaces the *entire string* of *from_string* into *to_string*, rather than one at a time like TRANSLATE.

```
SELECT 'Complete IT Professional' as SAMPLE_TEXT,
REPLACE('Complete IT Professional', 'let', 'sip') as REPLACE_TEST,
TRANSLATE('Complete IT Professional', 'let', 'sip') as TRANSLATE_TEST
FROM dual;
```

Result:

SAMPLE_TEXT	REPLACE_TEST	TRANSLATE_TEST
Complete IT Professional	Compsipe IT Professional	Compsipi IT Profissionas

MERGE

<https://www.oracletutorial.com/oracle-basics/oracle-merge/>

- A command for copying/merging data from a source table to a target table, utilizing matching conditions
- Template:

```
MERGE INTO target_table
USING source_table
ON search_condition
  WHEN MATCHED THEN
    UPDATE SET col1 = value1, col2 = value2,...
    WHERE <update_condition>
    [DELETE WHERE <delete_condition>]
  WHEN NOT MATCHED THEN
    INSERT (col1,col2,...)
    values(value1,value2,...)
    WHERE <insert_condition>;
```

MERGE

/* Update PRODUCT table with the data from NEW_PRODUCT
When their names matches, add the quantity; else insert a new row
with the new product info */

```
MERGE INTO product prod
  USING new_product new_prod
  ON (prod.product_name = new_prod.product_name)
  WHEN MATCHED THEN
    UPDATE SET prod.stock_qty = prod.stock_qty + new_prod.stock_qty
  WHEN NOT MATCHED THEN
    INSERT (product_name, stock_qty)
    VALUES (new_prod.product_name, new_prod.qty);
```

MERGE

- Example source: <https://www.oracletutorial.com/oracle-basics/oracle-merge/>
- /* Merge Member data to Member_Staging tables
 - When member_ID matches, but if Fname, Lname, or Rank are different, then use member table data. When member_ID does not match, insert */

members				member_staging			
MEMBER_ID	FIRST_NAME	LAST_NAME	RANK	MEMBER_ID	FIRST_NAME	LAST_NAME	RANK
1	Abel	Wolf	Gold	1	Abel	Wolf	Silver
2	Clarita	Franco	Platinum	2	Clarita	Franco	Platinum
3	Darryl	Giles	Silver	3	Darryl	Giles	Bronze
4	Dorthea	Suarez	Silver	4	Dorthea	Gate	Gold
5	Katrina	Wheeler	Silver	5	Katrina	Wheeler	Silver
6	Lilian	Garza	Silver	6	Lilian	Stark	Silver
7	Ossie	Summers	Gold				
8	Paige	Mcfarland	Platinum				
9	Ronna	Britt	Platinum				
10	Tressie	Short	Bronze				

MERGE

<https://www.oracletutorial.com/oracle-basics/oracle-merge/>

```

MERGE INTO member_staging x
USING members y
ON (x.member_id = y.member_id)
WHEN MATCHED THEN
    UPDATE SET x.first_name = y.first_name,
               x.last_name = y.last_name,
               x.rank = y.rank
WHERE x.first_name <> y.first_name OR
      x.last_name <> y.last_name OR
      x.rank <> y.rank
WHEN NOT MATCHED THEN
    INSERT(x.member_id, x.first_name, x.last_name, x.rank)
    VALUES(y.member_id, y.first_name, y.last_name, y.rank);
  
```

SQL Topics

- **SQL Basics**
 - Basic DDL and DML
 - Simple commands on a Single table
 - Joins (inner join, outer join, natural joins, multi-table joins)
 - Subquery
- **SQL for Data Engineering**
 - Date manipulation and Date Arithmetic
 - String manipulation
 - Merge
- **SQL for Analytics**
 - Aggregation, grouping, and ordering
 - OLAP functions
 - Ranking, partition by, buckets
 - Creating new derived columns

Analytic Functions in SQL

- SQL has many analytic functions:
 - Aggregation
 - OLAP (Online Analytical Functions)
 - Ranking
 - **RANK() & DENSE_RANK ()**
 - Ranking within subgroups
 - **PARTITION BY**
 - Creating new derived columns
 - **WIDTH_BUCKET ()**
 - **CASE...WHEN**
 - **COALESCE ()**
 - **LAG() and LEAD()**
 - **CUME_DIST()**

Aggregate Functions

- Aggregation functions *operate on a single column* of a table and *return a single value*
- Five aggregate functions:
 - **COUNT**: returns the *number of values* in a specified column
 - **SUM**: returns the *sum of the values* in a specified column
 - **AVG**: returns the *average of the values* in a specified column
 - **MIN**: returns the *smallest value* in a specified column
 - **MAX**: returns the *largest value* in a specified column
- Where to use:
 - In the **SELECT** list
 - In the **HAVING** clause

Aggregation

- Example of five aggregation functions
SELECT **MAX**(bonus), **MIN**(bonus), **AVG**(bonus), **COUNT**(bonus),
SUM(bonus)
FROM Instructor
WHERE position = 'assistant';
- How many *different* course titles are there?
SELECT **COUNT**(**DISTINCT** title) AS count
FROM Course;

Aggregation

- Note: What's wrong with the following?

```
SELECT instructorID, bonus
FROM Instructor
WHERE bonus > AVG(bonus);
```

Note: use aggregate functions only in **SELECT** and **HAVING** clause

- Why?**
- Rewrite using a subquery


```
SELECT instructorID, bonus
FROM Instructor
WHERE bonus > (SELECT AVG(bonus)
                FROM Instructor);
```

Descriptive Statistics in SQL

- SELECT** cus_fname, cus_lname, cus_balance
AVG(cus_balance) mean,
MEDIAN (cus_balance) median,
STATS_MODE (cus_balance) mode,
STDDEV (cus_balance) "Standard Deviation"
FROM Product;

CUSTOMER	
PK	<u>cus_code</u>
	cus_lname
	cus_fname
	cus_initial
	cus_areacode
	cus_phone
	cus_balance

- Computing median over partition


```
SELECT cus_fname, cus_lname, cus_balance
MEDIAN (cus_balance) OVER (PARTITION BY cus_areacode) AS median
FROM Product;
```
- Computing SD over unique values


```
SELECT cus_fname, cus_lname, cus_balance
STDDEV (DISTINCT cus_balance) FROM Product;
```

Aggregation with GROUP BY

Find the total number of instructors in each department and the sum of their bonus, respectively

```
SELECT deptCode, COUNT(instructorID) AS count, SUM(bonus) AS sum
FROM Instructor
GROUP BY deptCode
ORDER BY deptCode;
```

•Result:

deptCode	count	sum
acct	2	1100
math	2	300

Instructor
-instructorID(PK)
-fName
-lName
-ssn
-deptCode
-deptName
-position
-bonus

•**GROUP BY** must include all non-aggregate function column names in the **SELECT** list.

What's wrong with the following query?

Suppose we have the Instructor table as follows.

•Find the deptCode and deptName, and the total number of instructors in each department and the sum of their bonus, respectively

```
SELECT deptCode, deptName,
       COUNT(instructorID) AS count,
       SUM(bonus) AS sum
FROM Instructor
GROUP BY deptCode
ORDER BY deptCode;
```

Instructor
-instructorID(PK)
-fName
-lName
-ssn
-deptCode
-deptName
-position
-bonus

Aggregation with GROUP BY and HAVING

Oracle SQL*Plus

File Edit Search Options Help

```
SQL> SELECT U_CODE, COUNT(DISTINCT (P_CODE)), AVG(P_PRICE)
2 FROM PRODUCT
3 GROUP BY U_CODE;
```

U_CODE	COUNT(DISTINCT(P_CODE))	AVG(P_PRICE)
21225	2	8.47
21231	1	8.45
21344	3	12.49
23119	2	41.97
24288	3	155.593333
25595	3	89.63
	2	10.135

7 rows selected.

```
SQL> SELECT U_CODE, COUNT(DISTINCT (P_CODE)), AVG(P_PRICE)
2 FROM PRODUCT
3 GROUP BY U_CODE
4 HAVING AVG(P_PRICE) < 10;
```

U_CODE	COUNT(DISTINCT(P_CODE))	AVG(P_PRICE)
21225	2	8.47
21231	1	8.45

PRODUCT	
PK	<u>p_code</u>
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code

Aggregation with GROUP BY and HAVING

- Find the number of invoice per year and per month for the last 3 years

```
SELECT EXTRACT(year FROM inv_date) "Year",
       EXTRACT(month FROM inv_date) "Month",
       COUNT(inv_date) "No. of Invoices"
FROM invoice
GROUP BY EXTRACT(year FROM inv_date),
         EXTRACT(month FROM inv_date)
HAVING EXTRACT(year FROM inv_date) IN (2017, 2016, 2015)
ORDER BY "No. of Invoices" DESC;
```

INVOICE	
PK	<u>inv_number</u>
FK1	cus_code
	inv_date

Nested Subquery and Aggregate Functions

- Find the product that has the oldest date


```
SELECT P_CODE, P_DESCRIPT
FROM   PRODUCT
WHERE  P_INDATE = (
      SELECT MIN (P_INDATE)
      FROM   PRODUCT);
```

PRODUCT	
PK	<u>p_code</u>
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code

This query is in the form of a nested query.

SET Operations


```
-- All customers and staff
SELECT first_name, last_name
FROM customer
UNION
SELECT first_name, last_name
FROM staff
```



UNION

$R \cup S$


```
-- customers that are actors
SELECT first_name, last_name
FROM customer
INTERSECT
SELECT first_name, last_name
FROM actors
```



INTERSECT

$R \cap S$

```
-- customers that are not staff
SELECT first_name, last_name
FROM customer
EXCEPT
SELECT first_name, last_name
FROM staff
```



MINUS

$R - S$

SQL OLAP Functions: ROLLUP

- Example: **GROUP BY ROLLUP** (V_Code, P_Code) produces the union of
 - GROUP BY V_Code, P_Code
 - GROUP BY V_Code
 - Grand total

```
SQL> SELECT  V_CODE, P_CODE, SUM(SALE_UNITS*SALE_PRICE) AS TOTSALES
2  FROM      DWDAWSALESFACT NATURAL JOIN DWPRODUCT NATURAL JOIN DWVENDOR
3  GROUP BY  ROLLUP (V_CODE, P_CODE)
4  ORDER BY  V_CODE, P_CODE;
```

59

Oracle SQL*Plus

```
File Edit Search Options Help
SQL> SELECT  V_CODE, P_CODE, SUM(SALE_UNITS*SALE_PRICE) AS TOTSALES
2  FROM      DWDAWSALESFACT NATURAL JOIN DWPRODUCT NATURAL JOIN DWVENDOR
3  GROUP BY  ROLLUP (V_CODE, P_CODE)
4  ORDER BY  V_CODE, P_CODE;
```

V_CODE	P_CODE	TOTSALES
21225	23109-H8	99.5
21225	PUC23DRT	199.58
21225	SM-18277	41.94
21225		341.02
21344	13-Q2/P2	239.84
21344	54778-2T	59.88
21344		299.72
23119	1546-002	79.9
23119		79.9
24288	2232/PTY	219.84
24288	89-WRE-Q	513.98
24288		733.82
25595	2238/QPD	77.9
25595	WR3/TT3	719.7
25595		797.6
		2252.06

16 rows selected.

Subtotals by V_CODE

Grand total for all P_CODE values

```

graph TD
    DWTIME --> DWDAWSALESFACT
    DWVENDOR --> DWDAWSALESFACT
    DWVENDOR --> DWPRODUCT
    DWDAWSALESFACT --> DWPRODUCT
  
```

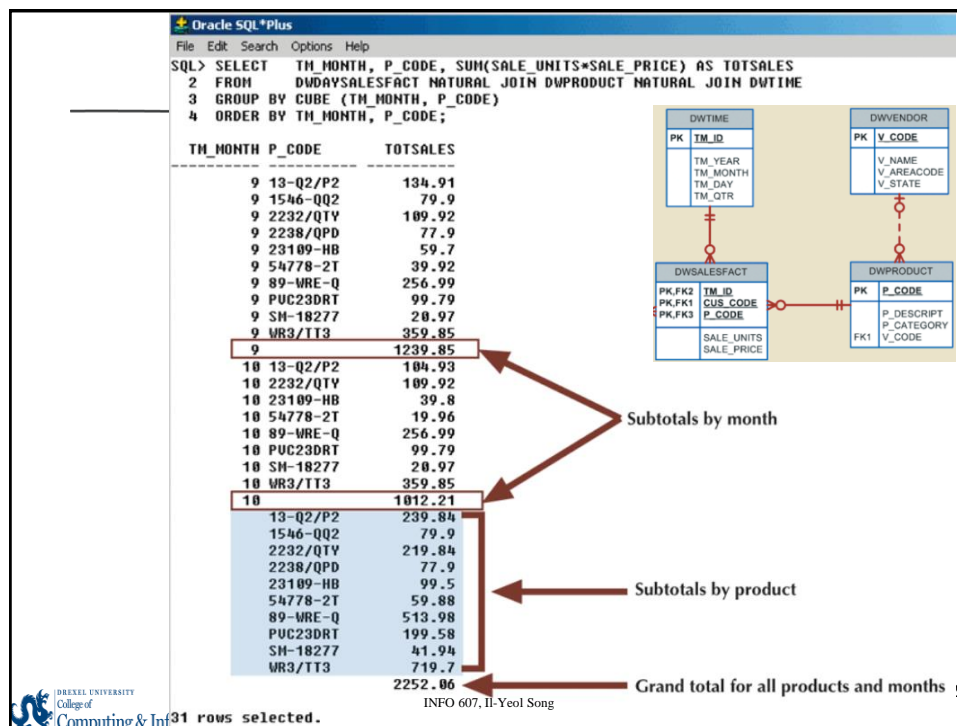
60

SQL OLAP Functions: Cube

- The CUBE extension
 - Enables you to get a subtotal for each column listed in the expression, in addition to a grand total for the last column listed
 - GROUP BY CUBE (TM_Month, P_CODE) is a union of:
 - GROUP BY (TM_Month, P_CODE)
 - GROUP BY (TM_Month)
 - GROUP BY (P_CODE)
 - Grand total

```
SQL> SELECT  TM_MONTH, P_CODE, SUM(SALE_UNITS*SALE_PRICE) AS TOTSALES
2 FROM      DWDAWSALESFACT NATURAL JOIN DWPRODUCT NATURAL JOIN DWTIME
3 GROUP BY CUBE (TM_MONTH, P_CODE)
4 ORDER BY TM_MONTH, P_CODE;
```

61



62

RANK and DENSE RANK Functions

- RANK() leave ranking gaps when there are multiple rows in the same rank (1, 2, 2, 4, 5...)
- DENSE_RANK() does **not** leave ranking gaps (1, 2, 2, 3, 4, 5)

```
SELECT p_code, p_descript, p_price,
       RANK() OVER
         (ORDER BY p_price NULLS LAST) AS Rank,
       DENSE_RANK ()
         (ORDER BY p_price NULLS LAST) AS Dense_rank
FROM   Product
ORDER BY p_price;
```

- DESC means descending order
- NULLS LAST means null values are smaller than non-null values
- You may use NULLS FIRST

PRODUCT	
PK	<u>p_code</u>
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code

RANK and DENSE RANK Functions

- Ranking applied to dates

```
SELECT p_code, p_descript, p_indate,
       RANK() OVER
         (ORDER BY p_indate NULLS LAST) AS Rank,
       DENSE_RANK ()
         (ORDER BY p_indate NULLS LAST) AS Dense_rank
FROM   Product
ORDER BY p_indate;
```

PRODUCT	
PK	<u>p_code</u>
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code

Top-N Queries

- **In Oracle 12c and after:**

```
SELECT      p_code, p_descript, p_price,
            RANK() OVER
              (ORDER BY p_price NULLS LAST) AS Rank,
FROM Product
ORDER BY p_price;
FETCH FIRST 10 ROWS ONLY;
```

- **Another option by percentage**

```
FETCH FIRST 20 PERCENT ROWS ONLY;
```

- **MySQL uses LIMIT clause**

```
LIMIT 10
```

PARTITION BY

- Partition the records by a column and then apply the function
- Ex: For each employee, show his/her dept# and the total# of employees working for the dept.

```
SELECT empno, deptno, COUNT(*)
      OVER (PARTITION BY deptno) DEPT_COUNT
FROM emp;
```

emp_no	dept_no	DEPT_COUNT
1	10	3
2	10	3
3	10	3 <- three because there are three "dept_no = 10" records
4	20	2
5	20	2 <- two because there are two "dept_no = 20" records

<https://stackoverflow.com/questions/561836/oracle-partition-by-keyword>

RANK with PARTITION BY

- Create ranks within each partition

SELECT employee_id, first_name, department_id, salary

RANK() OVER (PARTITION BY department_id

ORDER BY salary) rank

FROM employee;

EMPLOYEE_ID	FIRST_NAME	DEPARTMENT_ID	SALARY	RANK
1	119 Karen	30	2500	1
2	118 Guy	30	2600	2
3	117 Sigal	30	2800	3
4	116 Shelli	30	2900	4
5	115 Alexander	30	3100	5
6	114 Den	30	11000	6
7	144 Peter	50	2500	1
8	143 Randall	50	2600	2
9	142 Curtis	50	3100	3
10	141 Trena	50	3500	4
11	107 Diana	60	4200	1
12	105 David	60	4800	2

Windows Function: WIDTH_BUCKET

- This operation assigns a bucket number to the expression that it evaluates. It helps in generating equi-width histograms.
- Example: Show the last names, salaries and bucket numbers of all employees according to the following rule:

- Salary <30,000 -> 0 30,000-40,000 -> 1
- 40,000-50,000 -> 2 Salary >50,000 -> 3

SQL> SELECT lname, salary,

WIDTH_BUCKET(salary, 30000, 50000, 2) AS Bucket

FROM employee;

LNAME	SALARY	BUCKET
Smith	30000	1
Wong	40000	2
Zelaya	25000	0
Wallace	43000	2
Narayan	38000	1
English	25000	0
Jabbar	25000	0
Borg	55000	3
8 rows selected.		

CASE Statement

- Simple CASE Statements

```
SELECT      Fname, Lname,
            (CASE DNO
              WHEN 1      THEN 'Headquarters'
              WHEN 4      THEN 'Administration'
              WHEN 5      THEN 'Research'
              ELSE        'No department'
            END) AS Department
FROM Employee;
```

Output:

<u>Fname</u>	<u>Lname</u>	<u>Department</u>
John	Smith	Research
Franklin	Wong	Research
Alica	Zelaya	Administration

CASE Statement

- Searched CASE Statements

```
SELECT      Fname, Lname, Salary
            (CASE Salary
              WHEN Salary <= 25000 THEN 1500
              WHEN Salary > 25000 AND Salary < 50000 THEN 1000
              WHEN Salary > 50000 AND Salary < 100000 THEN 500
              ELSE 0
            END) AS "Bonus"
FROM Employee;
```

Output:

<u>Fname</u>	<u>Lname</u>	<u>Salary</u>	<u>Bonus</u>
John	Smith	30000	1000
Franklin	Wong	40000	1000
Alica	Zelaya	25000	1500

NULL and NVL Function

- Changing the display of null values
SQL> SET NULL N/A
All null values are displayed as N/A.
SQL> SET NULL -- to reset the null to blank
- Using **NVL** function to change the null value to another value (Only in Oracle)
SQL> SELECT Fname, Lname, NVL(CreditCode, 999)
FROM Cust;
 - If the value of CreditCode is not null, the value is displayed
 - If the value of CreditCode is null, 999 is displayed.
 - NVL(v1,v2): v1 and v2 must match by their data types.
 - But v1 can be any data type.

NULL and NVL Function

- Changing the display of null values
SQL> SET NULL N/A
All null values are displayed as N/A.
SQL> SET NULL -- to reset the null to blank
- Using **NVL** function to change the null value to another value (Only in Oracle)
SQL> SELECT Fname, Lname, NVL(CreditCode, 999)
FROM Cust;
 - If the value of CreditCode is not null, the value is displayed
 - If the value of CreditCode is null, 999 is displayed.
 - NVL(v1,v2): v1 and v2 must match by their data types.
 - But v1 can be any data type.

COALESCE Function

- A general form of NVL function-SQL 92 standard
- Provides a default value in case the column returns a NULL.

```
SELECT
    S.storecode,
    coalesce(s.date, '2021-05-10') as date,
    coalesce(s.salesqty,0) as salesqty,
    coalesce(s.salesrevenue, 0)as salesrevenue
FROM Products P;
```

- NVL uses only 2 arguments, while COALESCE can have multiple arguments, but stops at the first NOT NULL value
- Example:

Gives a 10% discount to all products with a list price. If there is no list price, then the sale price is the minimum price. If there is no minimum price, then the sale price is "5":

COALESCE Function

- Gives a 10% discount to all products with a list price. If there is no list price, then the sale price is the minimum price. If there is no minimum price, then the sale price is "5":

```
SELECT product_id, list_price, min_price,
    COALESCE(0.9*list_price, min_price, 5) Sale
FROM product_information
WHERE supplier_id = 102050
ORDER BY product_id, list_price, min_price, Sale;
```

PRODUCT_ID	LIST PRICE	MIN PRICE	Sale
1769	48		43.2
1770		73	73
2378	305	247	274.5
2382	850	731	765
3355			5

LAG(), LEAD()

- Lead() allows you to grab data from *a next row* without having to do a self join and put it into the current row
- Lag() allows you to grab data from *a previous row* and put it into the current row
- Syntax:

```
LEAD(scalar_expression ,offset [,default])
OVER (
    [PARTITION BY partition_expression, ... ]
    ORDER BY sort_expression [ASC | DESC], ...
```

LAG(), LEAD()

- Display employee joinDate and EndingDate in one row
- ```
SELECT empcode, empname, joindate,
 LEAD (joindate,1) OVER (ORDER BY joindate ASC)
 as EndDate
FROM employee;
```

| EMPCODE | EMPNAME  | JOININGDATE |
|---------|----------|-------------|
| 1       | Rajendra | 01-SEP-18   |
| 2       | Manoj    | 01-OCT-18   |
| 3       | Sonu     | 10-MAR-18   |
| 4       | Kashish  | 25-OCT-18   |
| 5       | Tim      | 01-DEC-18   |
| 6       | Akshita  | 01-NOV-18   |

| EMPCODE | EMPNAME  | JOININGDATE | ENDDATE   |
|---------|----------|-------------|-----------|
| 3       | Sonu     | 10-MAR-18   | 01-SEP-18 |
| 1       | Rajendra | 01-SEP-18   | 01-OCT-18 |
| 2       | Manoj    | 01-OCT-18   | 25-OCT-18 |
| 4       | Kashish  | 25-OCT-18   | 01-NOV-18 |
| 6       | Akshita  | 01-NOV-18   | 01-DEC-18 |
| 5       | Tim      | 01-DEC-18   | - null    |

## LAG(), LEAD()

- Given sales data below, display year, sales amount, and increase from the previous year

```
SELECT year, sales,
 Sale -LAG (sale, 1, NULL) OVER (ORDER BY year
ASC) AS increase-from-last-year
FROM sales_table;
```

| year | sale | year | sale | increase-from-last-year |
|------|------|------|------|-------------------------|
| 2017 | 100  | 2017 | 100  | NULL                    |
| 2018 | 125  | 2018 | 125  | 25                      |
| 2019 | 140  | 2019 | 140  | 15                      |
| 2020 | 175  | 2020 | 175  | 35                      |

## CUME\_DIST()

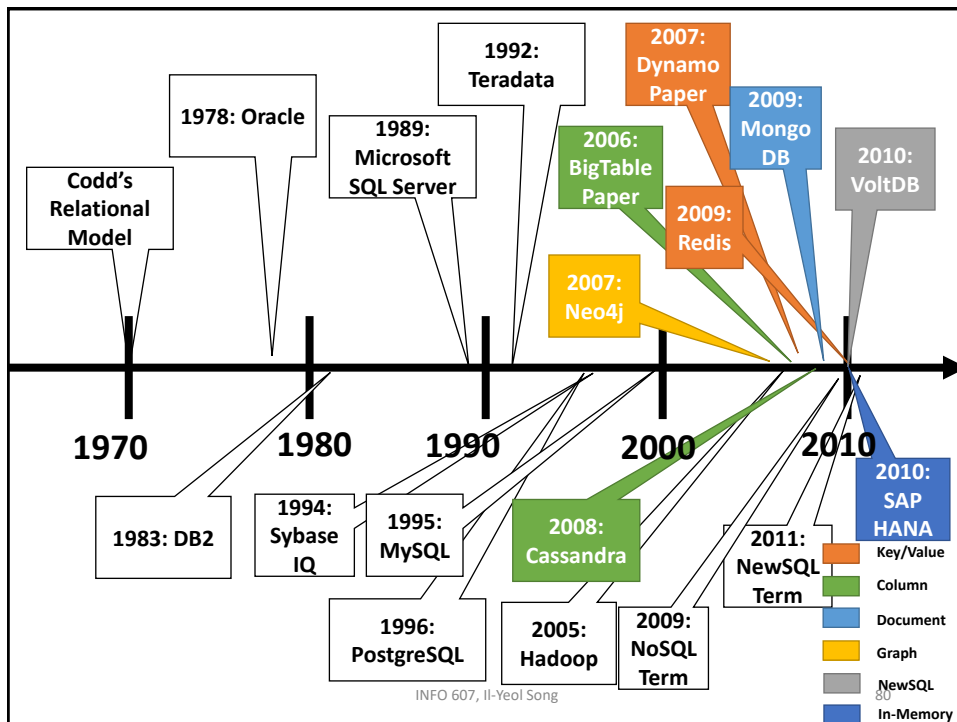
- CUME\_DIST() find the cumulative distribution of a value within a group of values, between 0 and 1.
  - Represents the number of rows that have values less than or equal to that row's value **divided by** the total number of rows

```
SELECT first_name, last_name, fees_paid,
 CUME_DIST() OVER (ORDER BY fees_paid) AS cume_dist_test
FROM student;
```

| FIRST_NAME | LAST_NAME | FEES_PAID | CUME_DIST_TEST |
|------------|-----------|-----------|----------------|
| Julie      | Armstrong | 0         | 0.07142857143  |
| Mark       | Anderson  | 45        | 0.1428571429   |
| Steven     | Webber    | 80        | 0.2142857143   |
| John       | Smith     | 100       | 0.4285714286   |
| Robert     | Pickering | 100       | 0.4285714286   |
| Mary       | Taylor    | 100       | 0.4285714286   |
| Susan      | Johnson   | 150       | 0.5714285714   |
| Tanya      | Hall      | 150       | 0.5714285714   |
| Jarrad     | Winston   | 300       | 0.6428571429   |
| Tom        | Capper    | 320       | 0.7142857143   |
| Andrew     | Cooper    | 400       | 0.7857142857   |
| Mark       | Holloway  | 410       | 0.8571428571   |
| John       | Rogers    | 700       | 0.9285714286   |
| Michelle   | Randall   | (null)    |                |

## Other Advanced SQL

- Other SQL functions and commands
- Views and Materialized views
- Indexes
- SQL with JSON
- SQL with Python
- PL/SQL
  - Stored functions
  - Stored Procedures
  - Triggers
  - Packages





## Major Features of NoSQL Databases (1)



- Manage **non-relational data**: Semi and non-structured data
- Stores **aggregated objects**:
  - Stores all the related data together (denormalized objects)
    - No complex relationships
  - Could cause redundant data and update anomaly
  - Write-once-read-many applications
- **Schema-on-read, mostly**
  - Some column stores need to define the schema in advance, but columns can be added dynamically
- Built on **the scale-out architecture**
  - Parallel processing, Scalability, fault tolerance, High availability (HA)
  - Support **automatic sharding**
    - Data are automatically distributed to nodes based on some fields

INFO 607, Il-Yeol Song

81

81

## Major Features of NoSQL Databases (2)



- Soften the ACID property:
  - **BASE**: Basically **Available**, Soft state, **Eventual consistency**
    - **NoSQLs focus on scalability and availability**
  - The data (and its replicas) will become consistent at some point in time after each transaction
  - Supports high throughput and low latency
- Open-sourced
- Four types
  - Key-value stores, Wide Column stores, Document stores
  - Graph databases
- Designed for **real-time OLTP system for non-uniform, big data.**

INFO 607, Il-Yeol Song

82

82

## Rankings of DB Systems by Popularity

- NoSQL is moving up.
  - DB-Engines Ranking *by popularity* (<https://db-engines.com/en/ranking> May, 2022): From Websites, Google trends, social networks, job Ads, etc.

| Rank     |          |          | DBMS                         | Database Model             |
|----------|----------|----------|------------------------------|----------------------------|
| May 2022 | Apr 2022 | May 2021 |                              |                            |
| 1.       | 1.       | 1.       | Oracle                       | Relational, Multi-model    |
| 2.       | 2.       | 2.       | MySQL                        | Relational, Multi-model    |
| 3.       | 3.       | 3.       | Microsoft SQL Server         | Relational, Multi-model    |
| 4.       | 4.       | 4.       | PostgreSQL                   | Relational, Multi-model    |
| 5.       | 5.       | 5.       | MongoDB                      | Document, Multi-model      |
| 6.       | 6.       | 7.       | Redis                        | Key-value, Multi-model     |
| 7.       | 8.       | 6.       | IBM Db2                      | Relational, Multi-model    |
| 8.       | 7.       | 8.       | Elasticsearch                | Search engine, Multi-model |
| 9.       | 9.       | 10.      | Microsoft Access             | Relational                 |
| 10.      | 10.      | 9.       | SQLite                       | Relational                 |
| 11.      | 11.      | 11.      | Cassandra                    | Wide column                |
| 12.      | 12.      | 12.      | MariaDB                      | Relational, Multi-model    |
| 13.      | 13.      | 13.      | Splunk                       | Search engine              |
| 14.      | 14.      | 27.      | Snowflake                    | Relational                 |
| 15.      | 15.      | 15.      | Microsoft Azure SQL Database | Relational, Multi-model    |
| 16.      | 16.      | 16.      | Amazon DynamoDB              | Multi-model                |
| 17.      | 17.      | 14.      | Hive                         | Relational                 |
| 18.      | 18.      | 17.      | Teradata                     | Relational, Multi-model    |
| 19.      | 19.      | 19.      | Neo4j                        | Graph                      |

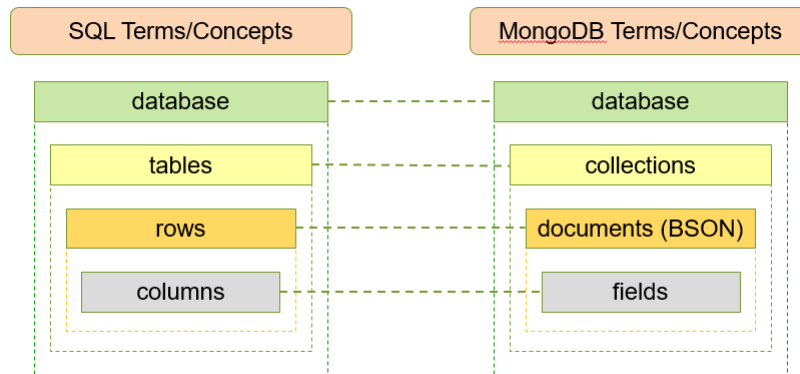
## MongoDB Shell

The following Mongo Shell provides a cloud based MongoDB engine. You can create database, insert/update records, queries, etc.

<https://docs.mongodb.com/manual/tutorial/query-embedded-documents/>

- The shell resets all data stored once closed.
- Create commands in Notepad and then copy and paste into the shell
- Be careful about quotes
- Up arrow display the previous command

## RDBMS and MongoDB



## JSON

- JavaScript Object Notation (JSON) – data interchange format used to represent data as a logical object.
- Every JSON document requires an object ID.
- **Object** is enclosed by a pair of **curly brackets { K:V }**
- **Array** is enclosed by a pair of **square brackets [ ]**
- **Value could be a “string”, an object, an array of values, a set of objects.**

- Example:

```
{
 _id: 101,
 title: "Database Systems",
 author: ["Coronel", "Morris"],
 publisher: {
 name: "Cengage",
 street: "500 Topbooks Avenue",
 city: "Boston",
 state: "MA"
 }
}
```

## CRUD Example

```
> db.user.insertOne({
 first: "John",
 last : "Doe",
 age: 39
})
```

```
> db.user.find ()
{
 "_id" : ObjectId("51..."),
 "first" : "John",
 "last" : "Doe",
 "age" : 39
}
```

```
> db.user.updateOne(
 { "_id" :
 ObjectId("51...") },
 { $set: {
 age: 40,
 salary: 7000 }
 }
)
```

```
> db.user.remove({
 "first": /^J/
})
```

## CRUD: Creating/Inserting documents

- Example:- **db.<<collectionName>>.insertOne({document})**

```
db.users.insertOne(← collection
{
 name: "sue", ← field: value
 age: 26, ← field: value
 status: "pending" ← field: value
} } document
)
```

```
>>> db.emp.insert({id: 1, name: "Song"})
WriteResult({ "nInserted" : 1 })
>>> |
```

## CRUD: Inserting Multiple documents


- Use the method – `insertMany([ {d1}, {d2}.., {dn} ] )`
- Example:-

```
db.inventory.insertMany([
 { item: "journal", qty: 25, tags: ["blank", "red"], size: { h: 14, w: 21, uom: "cm" } },
 { item: "mat", qty: 85, tags: ["gray"], size: { h: 27.9, w: 35.5, uom: "cm" } },
 { item: "mousepad", qty: 25, tags: ["gel", "blue"], size: { h: 19, w: 22.85, uom: "cm" } }
])
```

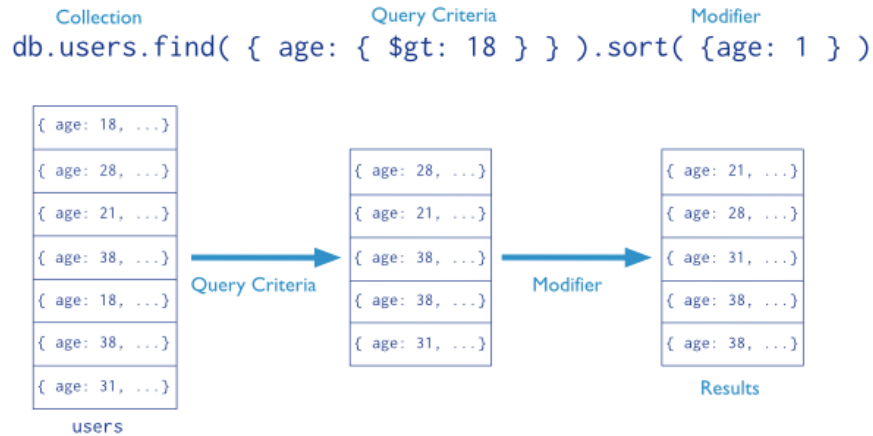
## CRUD: Find() with Query and Projections

Collection      Query Criteria      Projection

`db.users.find( { age: 18 }, { name: 1, _id: 0 } )`



## CRUD: find() with Query Criteria



## CRUD: UPDATE

- Update *all the documents* whose age is greater than 25

```
UPDATE people
SET status = "C"
WHERE age > 25
```

```
db.people.updateMany(
 { age: { $gt: 25 } },
 { $set: { status: "C" } }
)
```

- Update *the first document* whose age is greater than 25

```
db.people.updateOne(
 { age: { $gt: 25 } },
 { $set: { status: "C" } }
)
```

## CRUD: DELETE

```
DELETE FROM users
WHERE status = "D"
```

- Remove *the first* document from the users collection where the status field equals "D":

```
db.users.deleteOne({ status: "D" })
```

- Remove *all the documents* from the users collection where the status field equals "D":

```
db.users.deleteMany({ status : "D" })
```

- Remove *all the documents* from the users collection:

```
db.users.deleteMany({ })
```

## MongoDB: Other Topics

- Pattern Matching
- Array manipulation
- Aggregation framework
- Indexing
- Document design
- PyMongo
- Architecture
- Replication
- Sharding

## Good Resources

- **SQL Tutorial - Full Database Course for Beginners**, freeCodeCamp.org, <https://www.youtube.com/watch?v=HXV3zeQKqGY>
- **SQL Tutorial**, <https://www.w3schools.com/sql/>
- **Mongo DB Basics**, [Jyoti Reddy](#), Data Engineer at Krones
- <https://www.kdnuggets.com/2019/06/mongo-db-basics.html>
- **MongoDB Tutorial for Beginners: Learn Basics in 7 Days**, [David Taylor](#) Updated April 23, 2022. <https://www.guru99.com/mongodb-tutorials.html>

## Question?

